Sagemaker

Session-5

Table of Contents

I’ve completed the **pre-class** content?

SageMaker

Cloud Computing Market Size

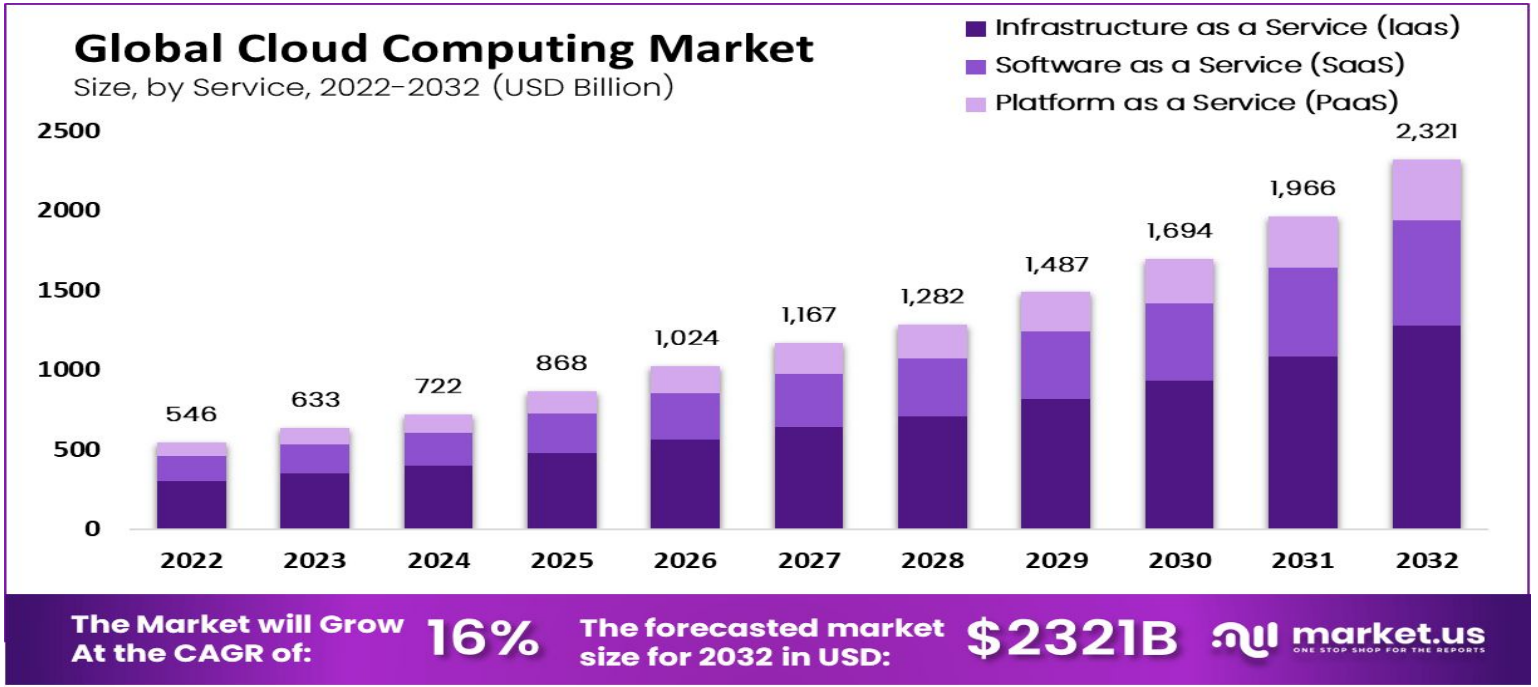

- The rising popularity of the latest novel technologies like **artificial intelligence** and **machine learning** and its rapid adoption in the cloud computing is empowering the growth of the global cloud computing market.

SageMaker

Alex Maese Juarez

What is SageMaker

Amazon **SageMaker** is a **platform** service that simplifies the process of **building, training, and deploying ML models** by providing everything organizations need to connect to their training data, select and optimize the **best algorithm and framework**, and **deploy their model** on auto-scaling clusters of Amazon **EC2**.

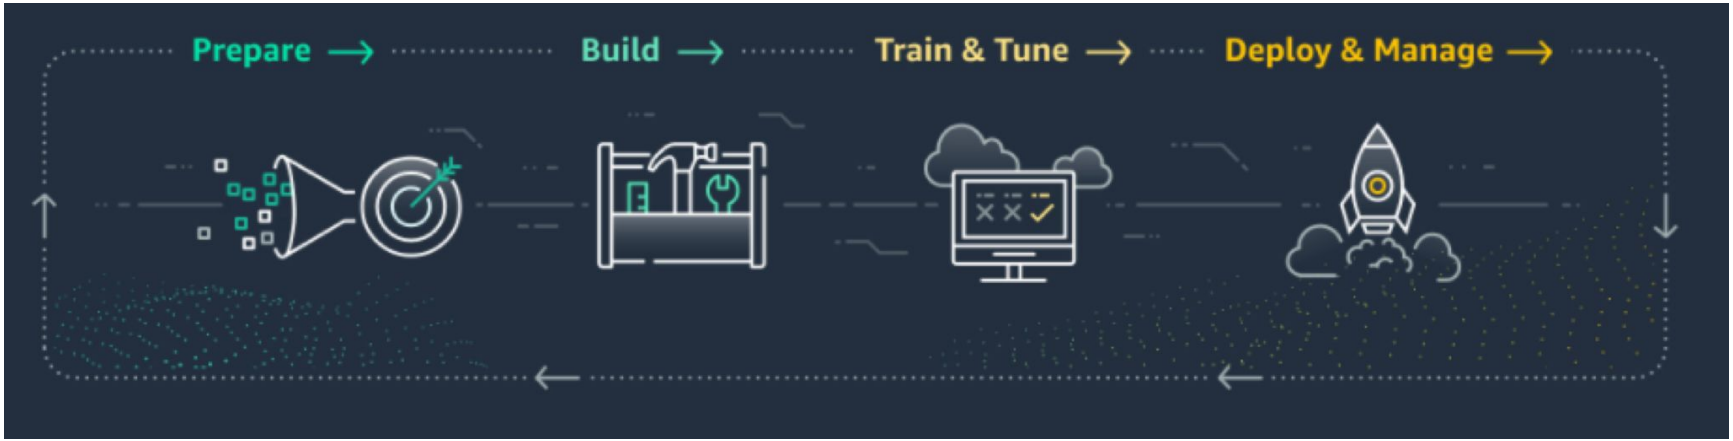

WHY Amazon SageMaker?

- Accelerating machine learning innovation through security
- Security features from Amazon SageMaker and the AWS Cloud can help organizations go from idea to production faster
- https://amer.resources.awscloud.com/ai-ml/accelerating-machine-learning-innovation-through-security

ML with AWS, by the numbers

- Reduce training time by 50%
- Provide 90% scaling efficiency
- Deliver 3x faster network throughput
- Improve price and performance by 25%
- 91% of cloud-based PyTorch runs on AWS
- 92% of cloud-based TensorFlow runs on AWS

ML process

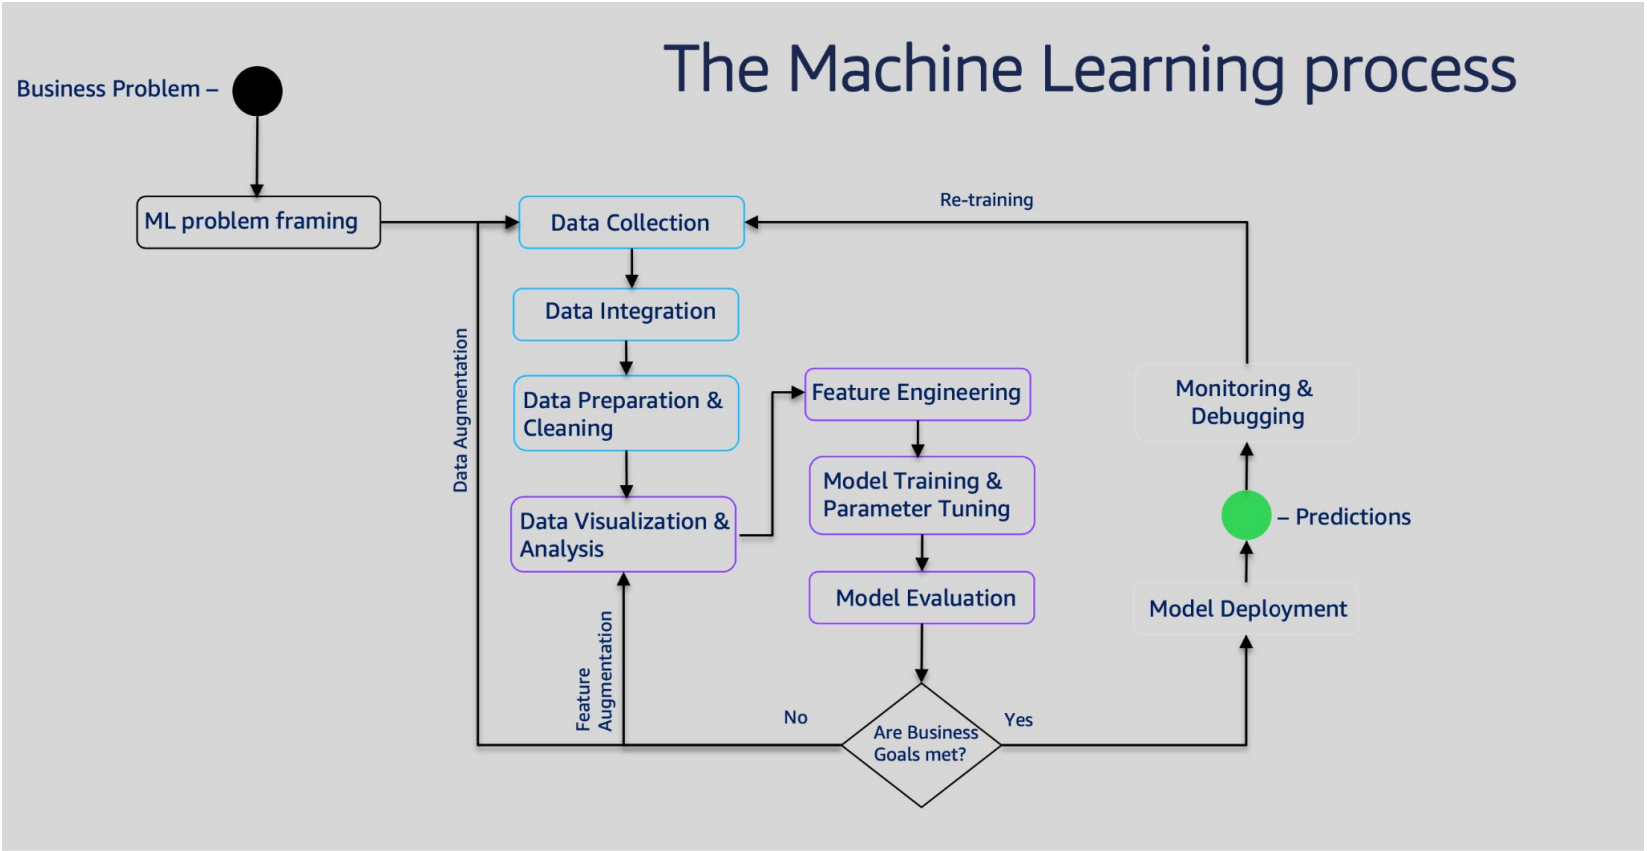
The Machine Learning process
Business Problem –
ML problem framing
Data Collection
Re-training
Data Integration
Data Augmentation
Data Preparation & Cleaning
Feature Engineering
Monitoring & Debugging
Model Training & Parameter Tuning
– Predictions
Data Visualization & Analysis
Model Evaluation
Model Deployment
Feature Augmentation
No
Are Business Goals met?
Yes

SageMaker Free Tier (2 months)

Two months free tier – starts from the first month you create a SageMaker resource

Development – 250 Hours/Month t2.medium or t3.medium

Train – 50 Hours/Month m4.xlarge or m5.xlarge

Deploy – 125 Hours/Month m4.xlarge or m5.xlarge

Model Deployment Architecture

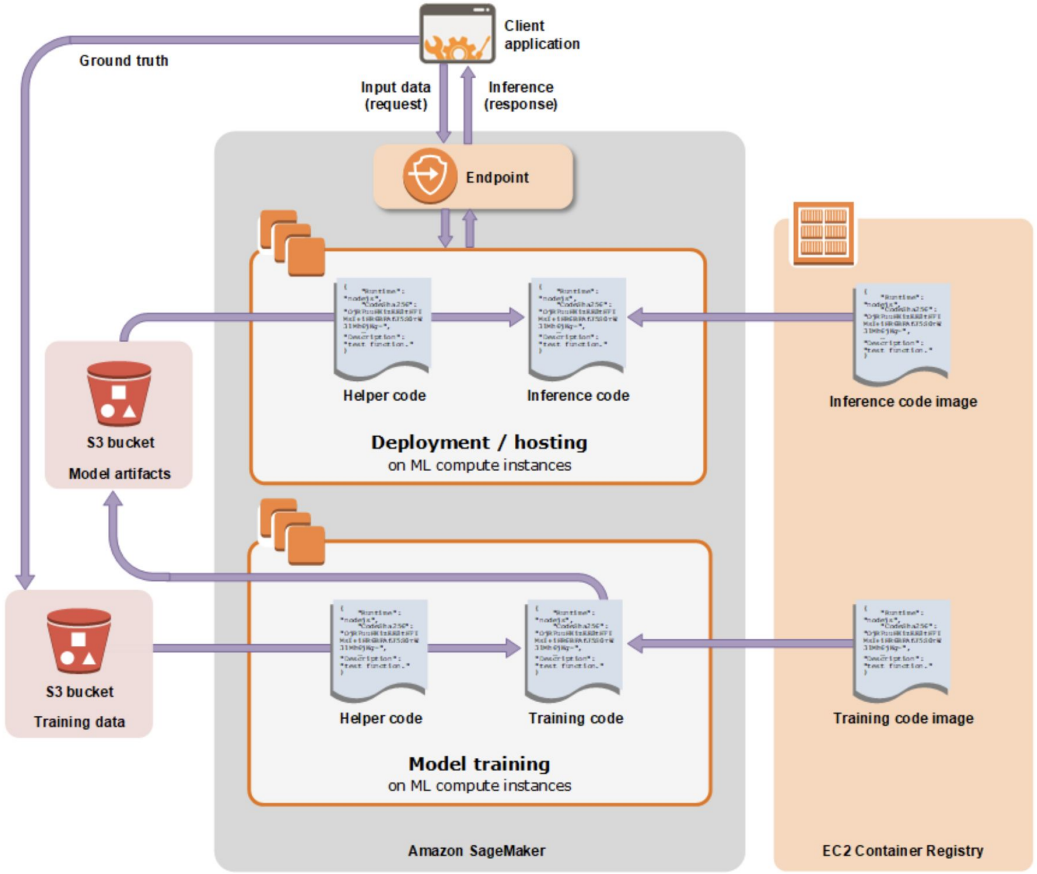

SageMaker Process

- Get your data

- Explore and refine models in a single Notebook Instance

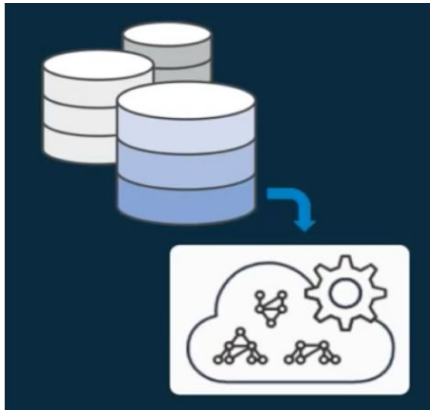

- Train on the full dataset in a cluster of GPU instances..

- Deploy to production

Built-in algorithms

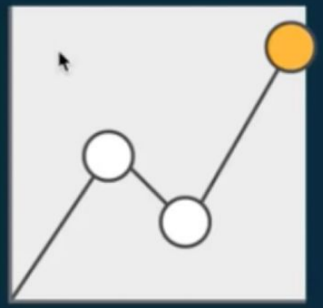

XGBoost, FM, Linear, k-NN, and Forecasting for supervised learning

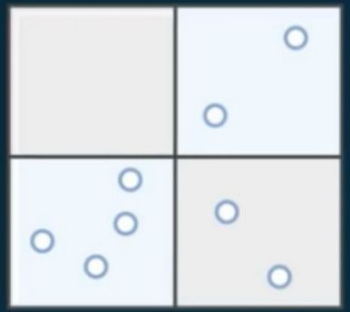

k-Means, PCA, and Random Cut Forest for unsupervised learning

Image classification and object detection for computer vision

LDA, Neural Topic Model, Seq2seq, and Word2Vec for text and NLP

Built-in algorithms-Supervised

- Linear Learner: regression
- K-Nearest Neighbors: non-parametric regression and classification
- XGBoost: regression, classification
- Factorization Machines: regression, classification, recommendation
- Semantic Segmentation: Deep Learning
- Image Classification: Deep Learning (ResNet)
- Object Detection (SSD): Deep Learning
- (VGG or ResNet)
- Sequence to Sequence: machine translation, speech to text and more
- DeepAR: time-series forecasting (RNN)

Built-in algorithms-Unsupervised

- K-Means: clustering
- Principal Component Analysis: dimensionality reduction
- Random Cut Forest: anomaly detection
- Object2Vec: general-purpose embedding
- Neural Topic Model: topic modeling
- Latent Dirichlet Allocation: topic modeling (mostly)
- Blazing Text: GPU-based Word2 Vec, and text classification
- IP Insights: usage patterns for IP addresses

Bring your own algorithm

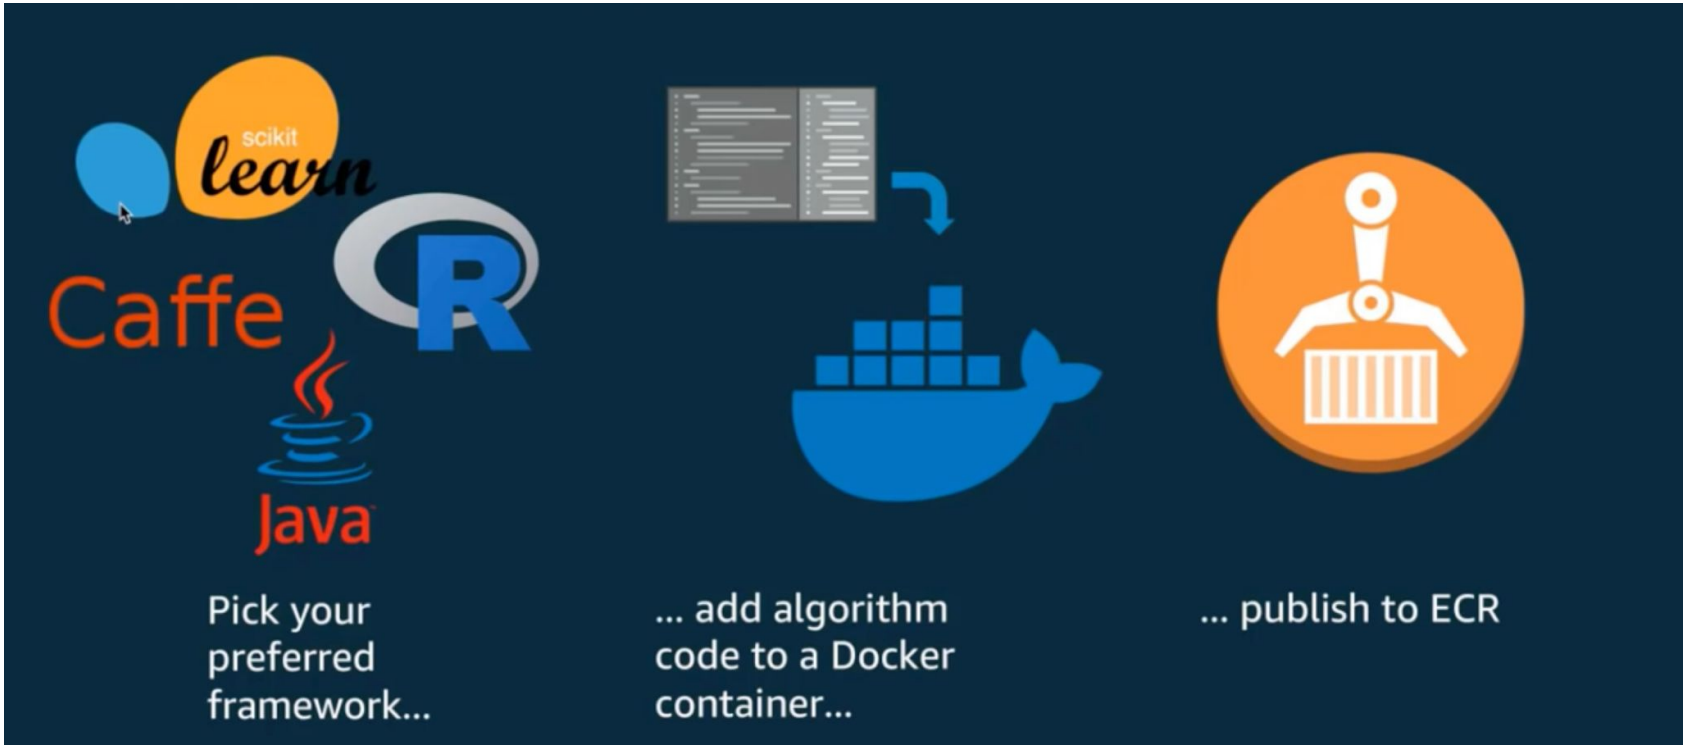

Let’s jump to the AWS SageMaker Service!

S3 BUCKET and OBJECTS

- iden-ml-sagemaker/
  - bikerental/
    - train/
      - train.csv
    - validation/
      - validation.csv
    - test/
      - test.csv
    - model/
      - xgboost-bikerental-v1-2022-04-21-10-23-10-964/ (Job_name+datetime)
        - output
          - model.tar.gz

Amazon S3>Buckets>iden-ml-sagemaker>bikerental/output/xgbobikerental-v1-2022-04-21-10-23-10-964/output/model.tar.gz

Notebook Instance

An *Amazon SageMaker notebook instance* is a machine learning (ML) compute instance running the Jupyter Notebook App. SageMaker manages creating the instance and related resources.

Notebook Instance

Use Jupyter notebooks in your notebook instance

- To prepare and process data,
- Write code to train models,
- Deploy models to SageMaker hosting,
- Test or validate your models.

Notebook Instance

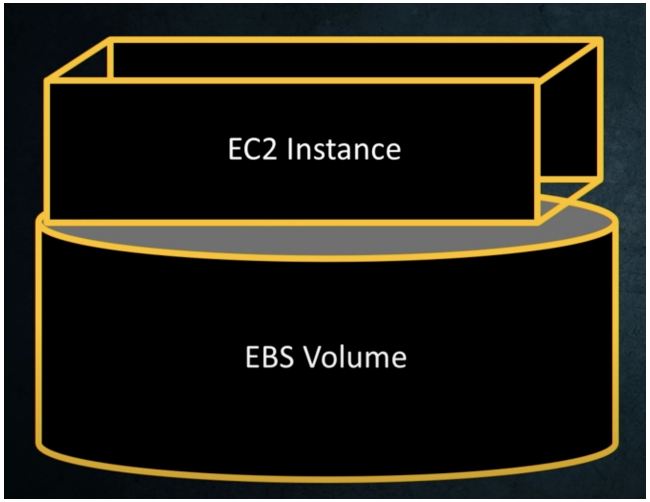

- Anaconda Packages,
- Tensorflow and Apache MXnet,
- Storage volume,
- Sample notebooks that contain complete code walkthroughs

You are using CSV formatted files to train on SageMaker's built-in XGBoost algorithm. SageMaker expects your training and validation to follow this convention:

- CSV must have column headers and target variable must be the last column
- CSV must have column headers with the target variable in the first column
- CSV must not have a column header record. Target variable must be the last column
- CSV must not have a column header record. Target variable must be the first column

How does SageMaker built-in know the target variable?

- For CSV training, the algorithm assumes that the target variable is in the first column and that the CSV does not have a header record.(Train,Validation)
- For CSV test, the algorithm assumes that CSV input does not have the label column.

https://docs.aws.amazon.com/sagemaker/latest/dg/xgboost.html

SageMaker Free Tier Usage

Amazon SageMaker capability	Free Tier usage per month for the first 2 months
Studio notebooks, and notebook instances	250 hours of ml.t3.medium instance on Studio notebooks OR 250 hours of ml.t2 medium instance or ml.t3.medium instance on notebook instances
RStudio on SageMaker	250 hours of ml.t3.medium instance on RSession app AND free ml.t3.medium instance for RStudioServerPro app
Data Wrangler	25 hours of ml.m5.4xlarge instance
Feature Store	10 million write units, 10 million read units, 25 GB storage (standard online store)
Training	50 hours of m4.xlarge or m5.xlarge instances
Amazon SageMaker with TensorBoard	300 hours of ml.r5.large instance
Real-Time Inference	125 hours of m4.xlarge or m5.xlarge instances
Serverless Inference	150,000 seconds of on-demand inference duration
Canvas	160 hours/month for session time
HyperPod	50 hours of m5.xlarge instance
	Free Tier usage per month for the first 6 months
Experiments	100,000 metric records ingested per month, 1 million metric records retrieved per month, and 100,000 metric records stored per month

https://aws.amazon.com/sagemaker/pricing/

SageMaker Billing Dashboard

Cost Management
- Cost Explorer
- Budgets
- Budgets Reports
- Savings Plans

Preferences
- Billing preferences
- Payment methods
- Consolidated billing
- Tax settings

▸ **EC2 Container Registry (ECR)**		**$0.00**
▸ **Elastic Compute Cloud**		**$0.00**
▸ **Key Management Service**		**$0.00**
▾ **SageMaker**		**$0.30**
▾ **US East (N. Virginia)**		**$0.30**
Amazon SageMaker CreateVolume-Gp2		$0.00
$0.00 for SageMaker Debugger Built-in Rule Volume	0.026 GB-Mo	$0.00
$0.14 per GB-Mo of Endpoint ML storage	0.002 GB-Mo	$0.00
$0.14 per GB-Mo of Notebook Instance ML storage	0.034 GB-Mo	$0.00
$0.14 per GB-Mo of Training Job ML storage	0.000278 GB-Mo	$0.00
Amazon SageMaker Invoke-Endpoint		$0.00
$0.016 per GB for Endpoint Data IN	0.001 GB	$0.00
$0.016 per GB for Endpoint Data OUT	0.000230 GB	$0.00
Amazon SageMaker RunInstance		$0.30
$0.0 for SageMaker Debugger Built-in Rule Instance	0.005 Hrs	$0.00
$0.00 for Host:ml.m4.xlarge per hour under monthly free tier	0.169 Hrs	$0.00
$0.00 for Notebk:ml.t2.medium per hour under monthly free tier	2.695 Hrs	$0.00
$0.00 for Notebk:ml.t3.medium per hour under monthly free tier	1.153 Hrs	$0.00
$0.23 per Notebook ml.m5.xlarge hour in US East (N. Virginia)	1.258 Hrs	$0.29
$0.478 per Training ml.c4.2xlarge hour in US East (N. Virginia)	0.016 Hrs	$0.01
Managed Spot Training Job ml.m5.xlarge hour in US East (N. Virginia)	0.003 Hrs	$0.00
▸ **US East (Ohio)**		**$0.00**
▸ **Secrets Manager**		**$0.00**
▸ **Simple Notification Service**		**$0.00**
▸ **Simple Storage Service**		**$0.00**

SageMaker Use Case Example

ZAPPOS

"We are…using analytics and machine learning solutions to personalize sizing and search results for individual users. AWS services (including Amazon SageMaker) allow (our) engineers to focus on improving performance and results rather than DevOps overhead."

Hyperparameters

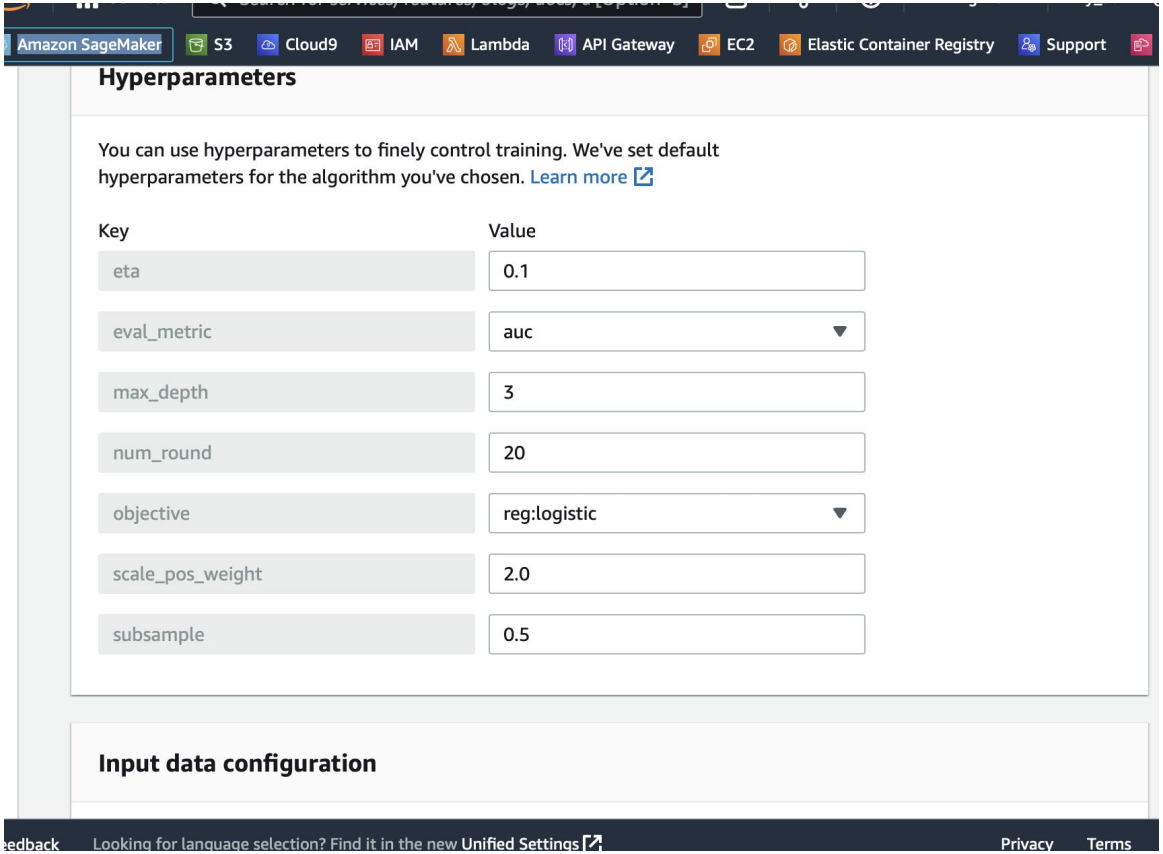

THANKS!

Any questions?

S3 Settings

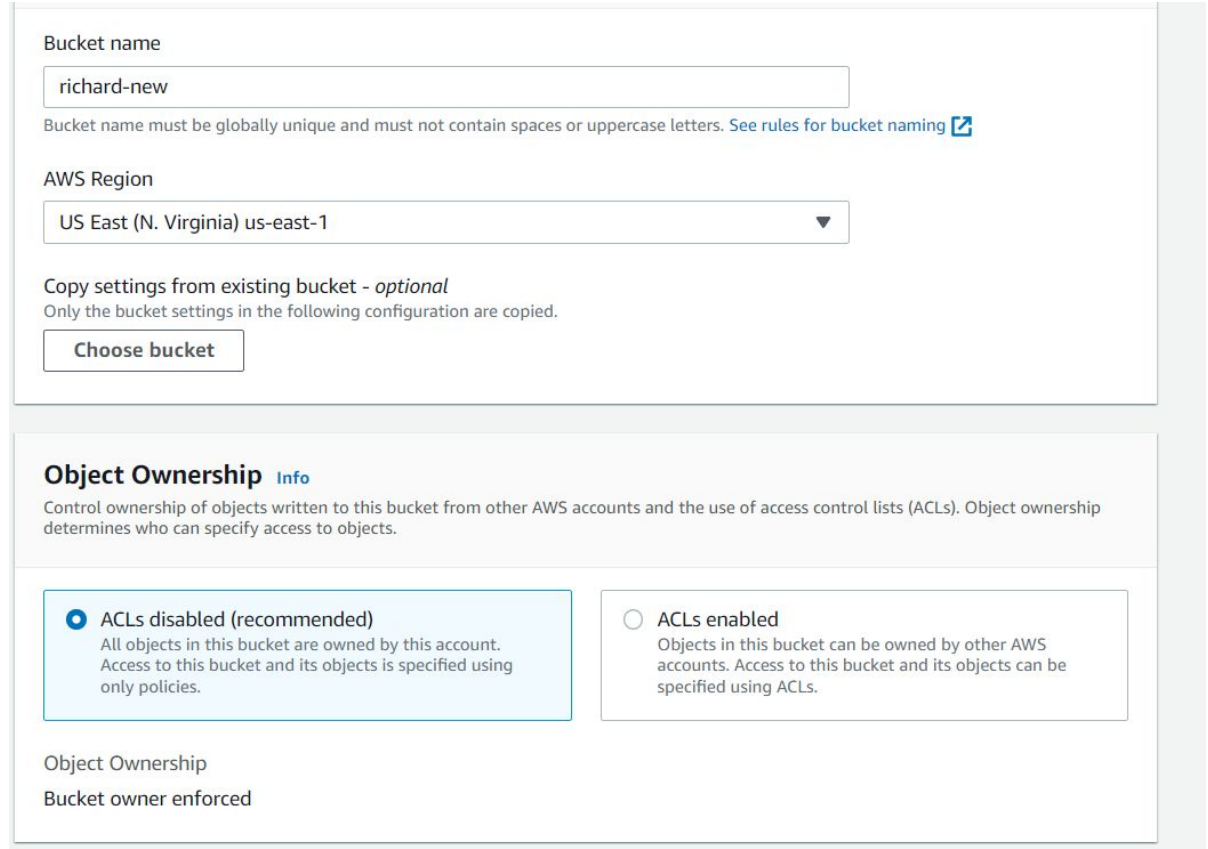
Bucket name
richard-new
Bucket name must be globally unique and must not contain spaces or uppercase letters. See rules for bucket naming
AWS Region
US East (N. Virginia) us-east-1
Copy settings from existing bucket - optional
Only the bucket settings in the following configuration are copied.
Choose bucket
Object Ownership Info
Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.
ACLs disabled (recommended)
All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.
ACLs enabled
Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.
Object Ownership
Bucket owner enforced

S3 Settings

Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. Learn more

- [x] **Block *all* public access**
  Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.
  - [x] **Block public access to buckets and objects granted through *new* access control lists (ACLs)**
    S3 will block public access permissions applied to newly added buckets or objects, and prevent the creation of new public access ACLs for existing buckets and objects. This setting doesn't change any existing permissions that allow public access to S3 resources using ACLs.
  - [x] **Block public access to buckets and objects granted through *any* access control lists (ACLs)**
    S3 will ignore all ACLs that grant public access to buckets and objects.
  - [x] **Block public access to buckets and objects granted through *new* public bucket or access point policies**
    S3 will block new bucket and access point policies that grant public access to buckets and objects. This setting doesn't change any existing policies that allow public access to S3 resources.
  - [x] **Block public and cross-account access to buckets and objects through *any* public bucket or access point policies**
    S3 will ignore public and cross-account access for buckets or access points with policies that grant public access to buckets and objects.

S3 Settings

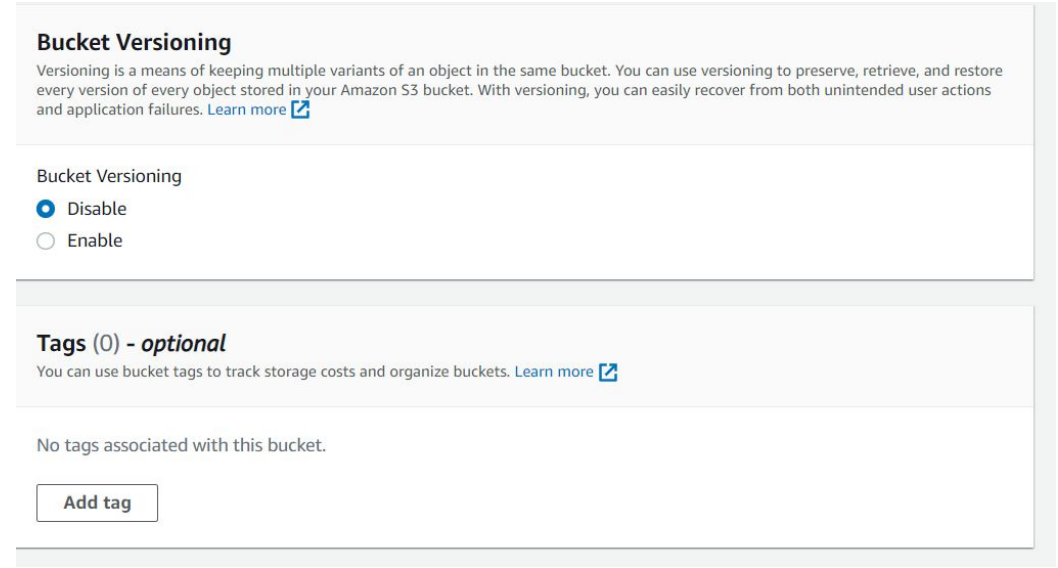
Bucket Versioning
Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. Learn more
Bucket Versioning
Disable
Enable
Tags (0) - optional
You can use bucket tags to track storage costs and organize buckets. Learn more
No tags associated with this bucket.
Add tag

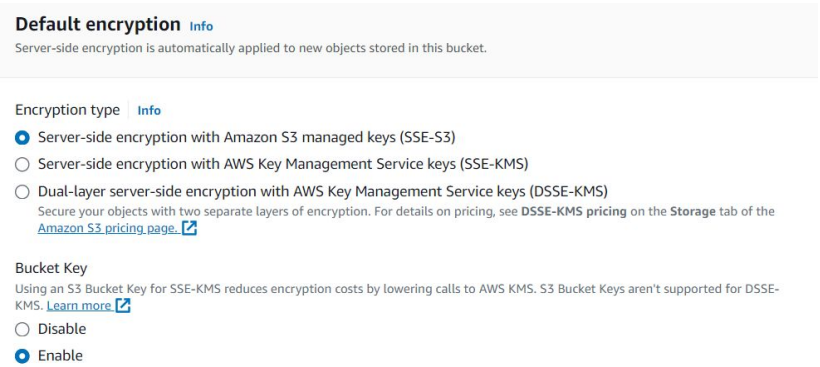
Default encryption Info
Server-side encryption is automatically applied to new objects stored in this bucket.
Encryption type Info
Server-side encryption with Amazon S3 managed keys (SSE-S3)
Server-side encryption with AWS Key Management Service keys (SSE-KMS)
Dual-layer server-side encryption with AWS Key Management Service keys (DSSE-KMS)
Secure your objects with two separate layers of encryption. For details on pricing, see DSSE-KMS pricing on the Storage tab of the Amazon S3 pricing page.
Bucket Key
Using an S3 Bucket Key for SSE-KMS reduces encryption costs by lowering calls to AWS KMS. S3 Bucket Keys aren't supported for DSSE-KMS. Learn more
Disable
Enable

Notebook Settings

Notebook instance settings

Notebook instance name

richard-sda-test

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Notebook instance type

ml.t3.medium

Elastic Inference Learn more

none

Platform identifier Learn more

Amazon Linux 2, Jupyter Lab 3

▶ Additional configuration

Notebook Settings

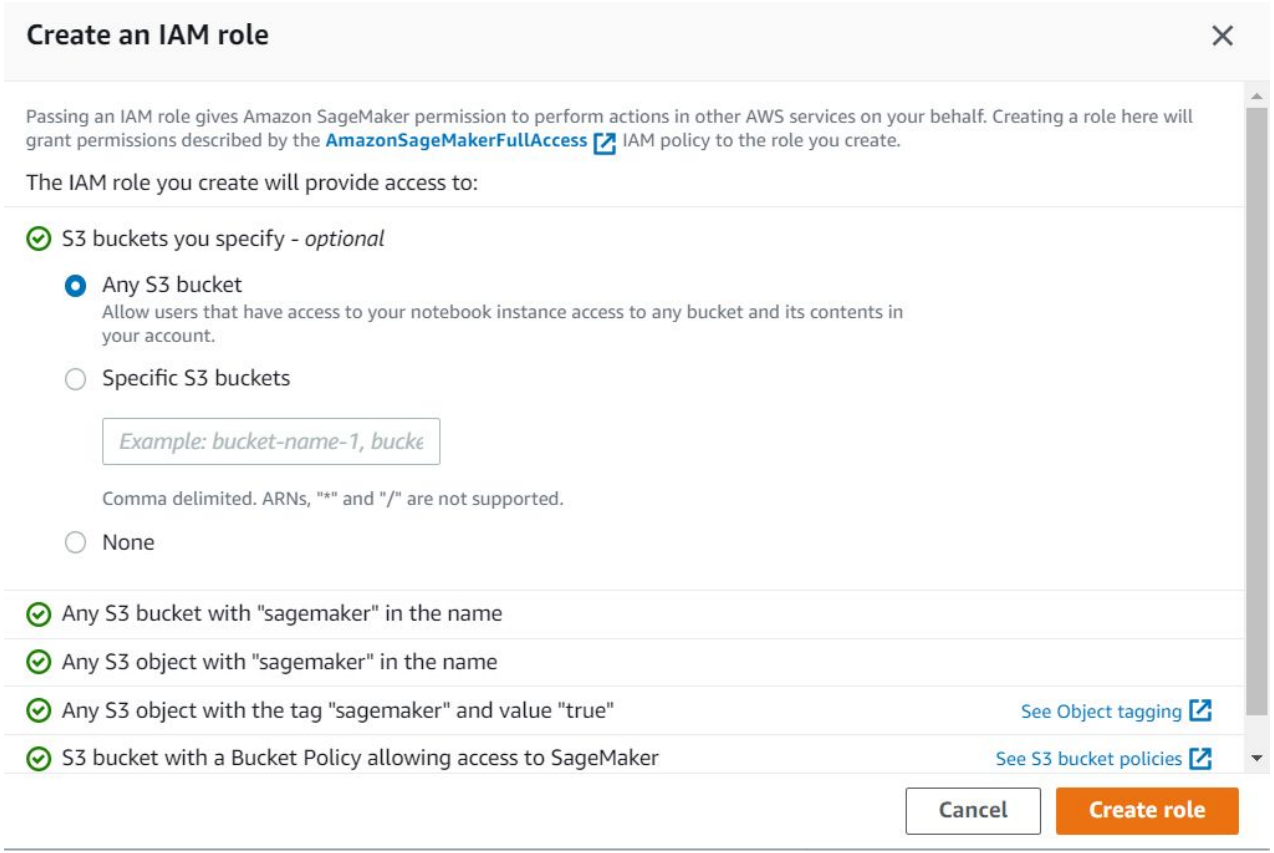

Notebook Settings

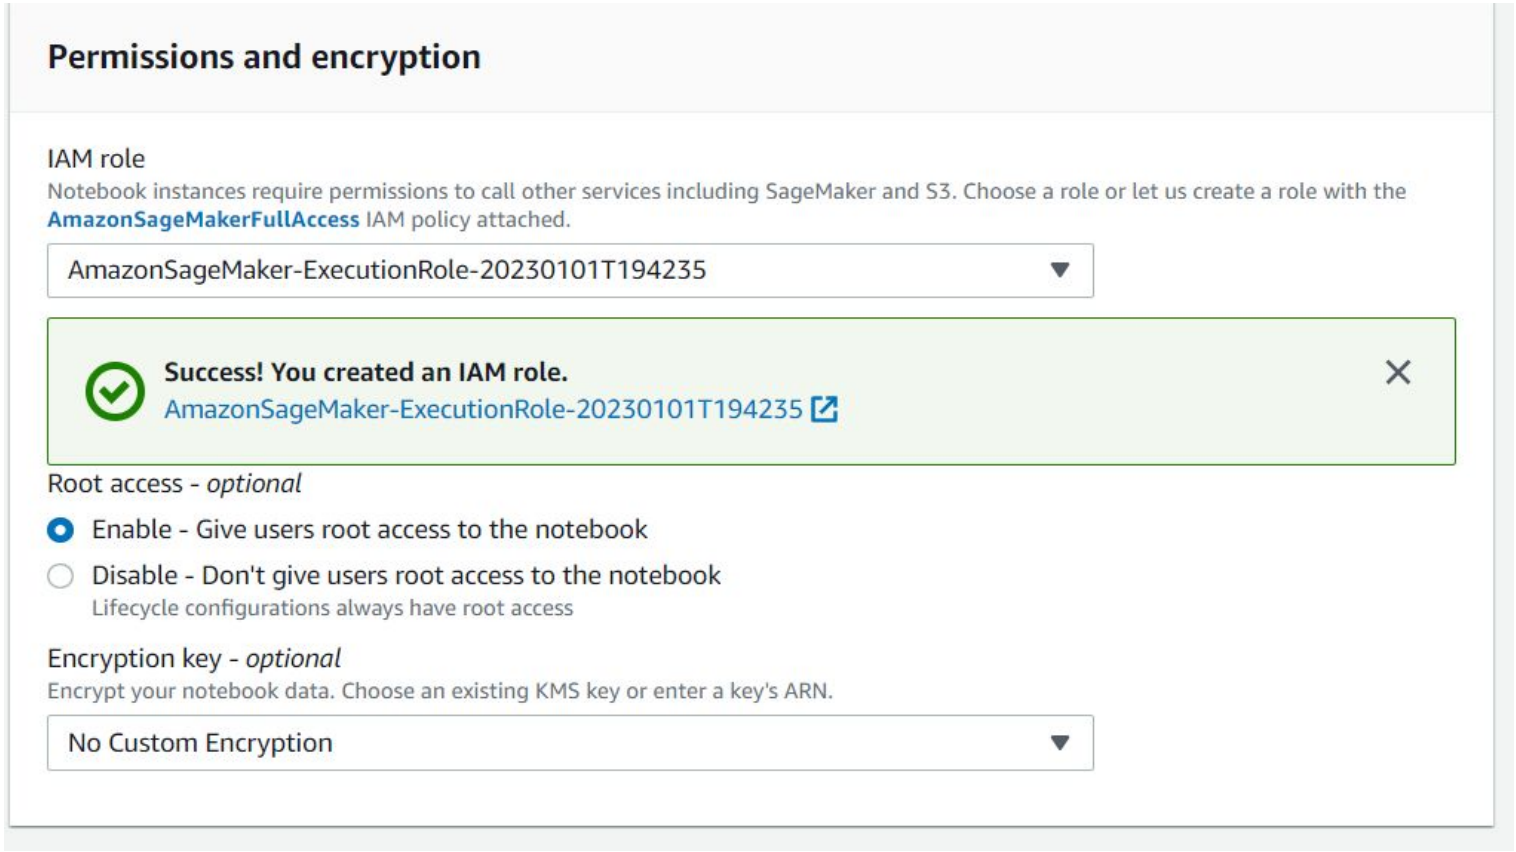

Notebook Settings

▶ **Network** - *optional*

▶ **Git repositories** - *optional*

▶ **Tags** - *optional*

Cancel **Create notebook instance**

Notebook Settings

Now open the notebook from Sagemaker and send the data S3 via notebook

After sending data come back to the sagemaker console and do the jobs:

Training Job Creation-Model Creation- End Point Creation

Train Job Settings

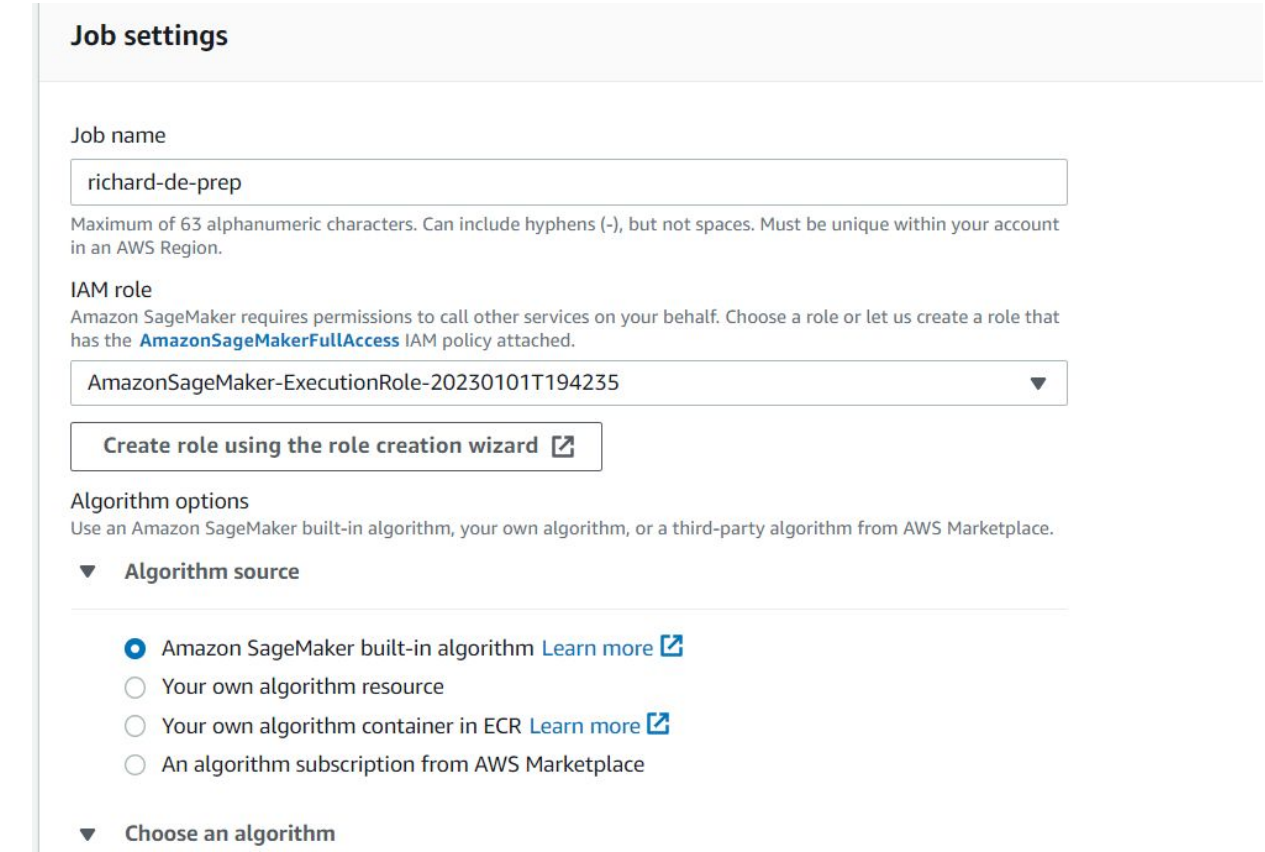

Train Job Settings

Choose an algorithm

Tabular - XGBoost : v1.3

Container

The registry path where the training image is stored in Amazon ECR. **Learn more**

683313688378.dkr.ecr.us-east-1.amazonaws.com/sagemaker-xgboost:1.3-1

Input mode

You can provide your training data as a file or pipe.

File

Metrics

The algorithm you selected will publish the following metrics to CloudWatch metrics.

Metric name	Regex
train:rmse	.*\[[0-9]+\].*#011train-rmse:([-+]?[0-9]*\.?[0-9]+(?:[eE][-+]?[0-9]+)?).*
train:mae	.*\[[0-9]+\].*#011train-mae:([-+]?[0-9]*\.?[0-9]+(?:[eE][-+]?[0-9]+)?).*
train:logloss	.*\[[0-9]+\].*#011train-logloss:([-+]?[0-9]*\.?[0-9]+(?:[eE][-+]?[0-9]+)?).*
. .	.*\[[0-9]+\].*#011train-error:([-+]?[0-9]*\.?[0-9]+(?:[eE][-+]?[0-

Train Job Settings

Resource configuration

Instance type: ml.m4.xlarge

Instance count: 1

Additional storage volume per instance (GB): 1

Keep alive period

Use **SageMaker Training Managed Warm Pools**

minutes

Maximum runtime is 1 hour (60 minutes or 3600 seconds).

Encryption key - *optional*

Encrypt your data. Choose an existing KMS key or enter a key's ARN.

No Custom Encryption

Stopping condition

Maximum runtime

1800

seconds

Train Job Settings

```
xgb_model.set_hyperparameters(max_depth=4,
                              eta=0.3,
                              num_round=200,
                              objective = "reg:squarederror",
                              early_stopping_rounds=10)  #objective = "reg:squarederror", "reg:linear" for latest xgboost!
```

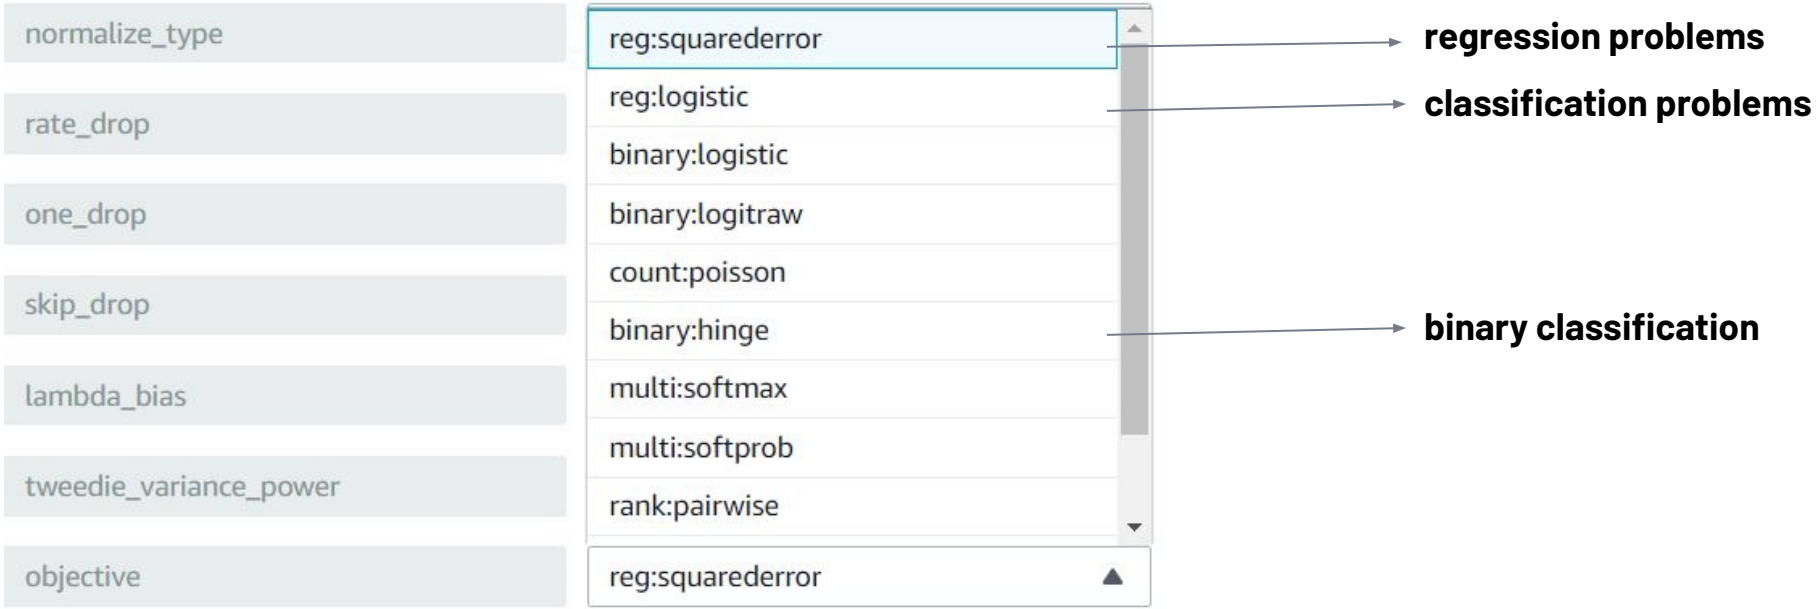

Train Job Settings

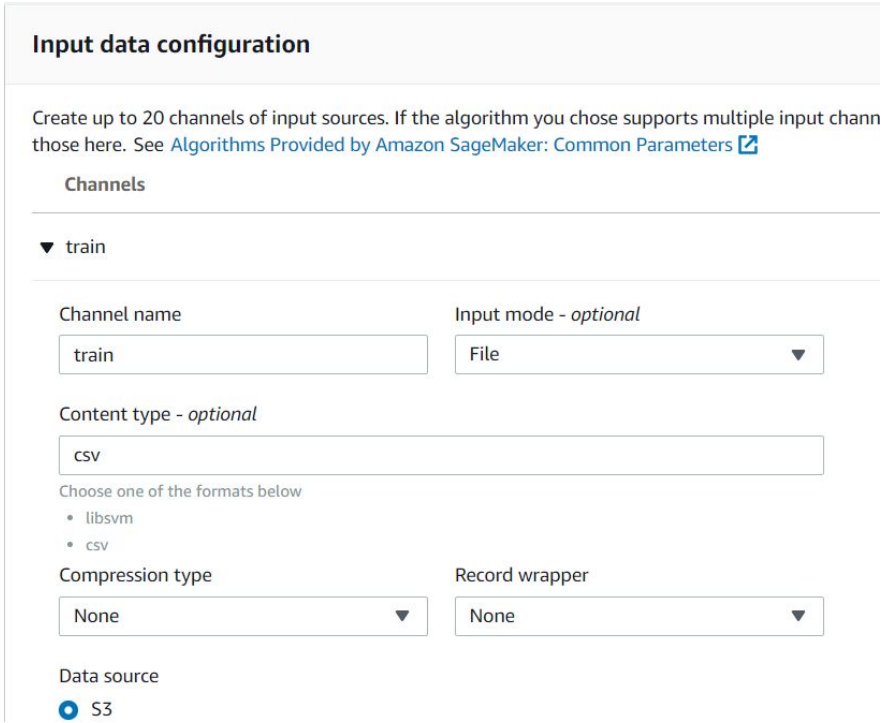

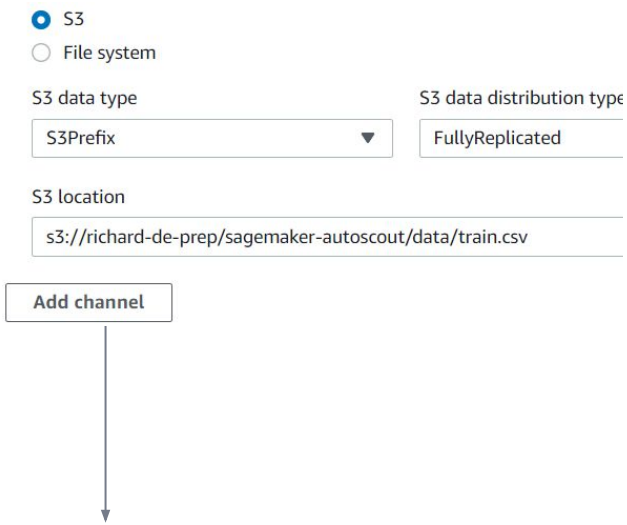

Don’t forget to add channel for validation data

Train Job Settings

▼ validation

Channel name

validation

Input mode - *optional*

File

Content type - *optional*

csv

Choose one of the formats below

- libsvm
- csv

Compression type

None

Record wrapper

None

Data source

- (•) S3
- ( ) File system

S3 data type

S3Prefix

S3 data distribution type

FullyReplicated

S3 location

s3://richard-de-prep/sagemaker-autoscout/data/validation.csv

Create Output Folder from S3

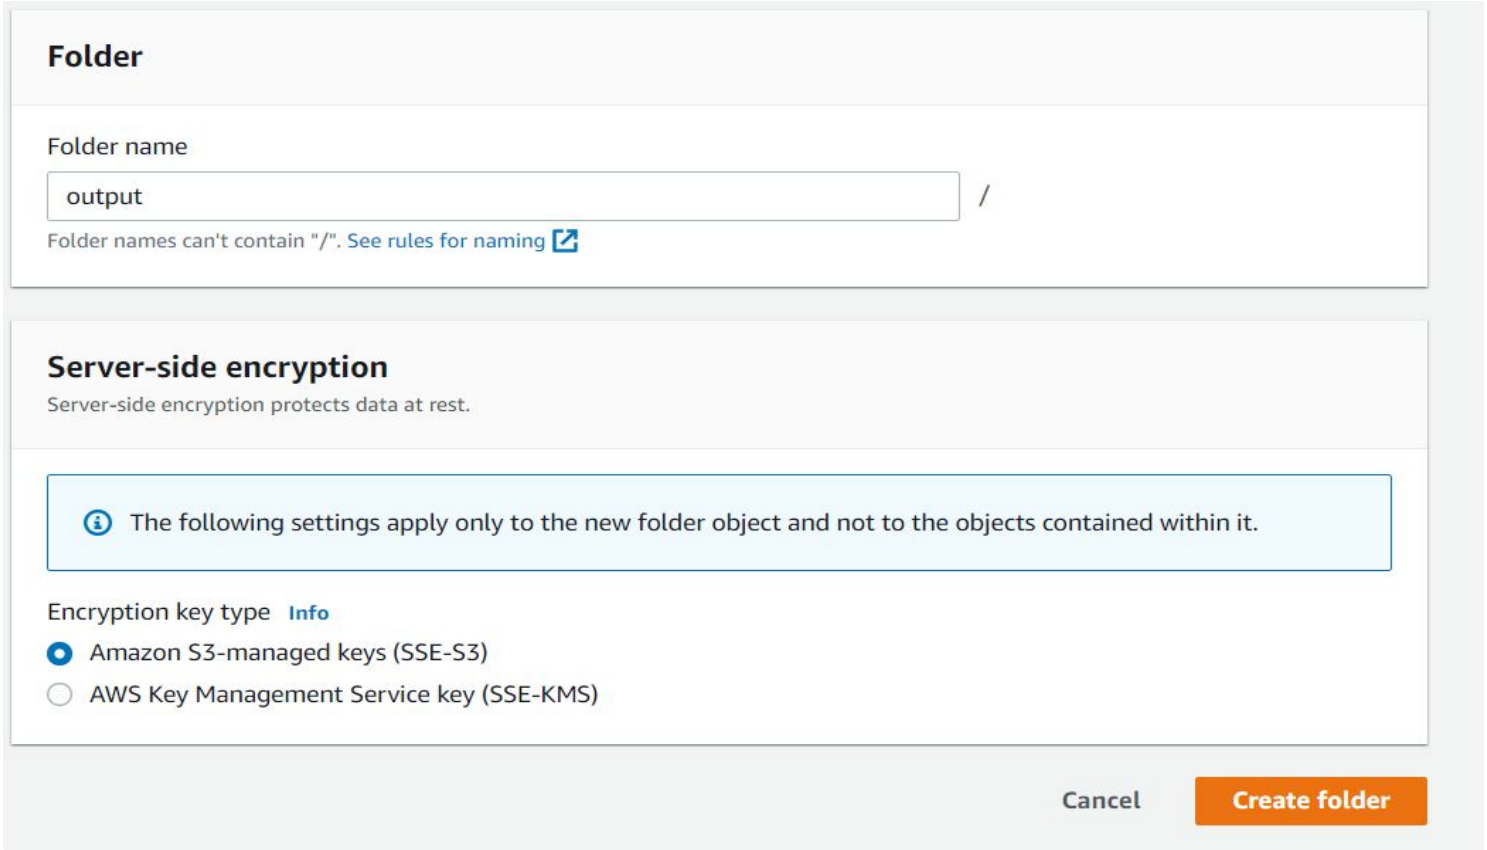

Train Job Settings

Output data configuration

S3 output path

s3://richard-de-prep/sagemaker-autoscout/output/

Encryption key - *optional*

If you want Amazon SageMaker to encrypt the output of your training job using your own AWS KMS encryption key instead of the default S3 service key, provide its ID or ARN.

Train Job Settings

Managed spot training

Enable managed spot training - *optional*

Save compute costs for jobs that have flexibility in start and end times. Amazon SageMaker will use spare capacity only to run this job. Learn more

Maximum wait time before job terminates *optional stopping condition*

At the end of this duration you will receive the complete or partial results of you managed spot training job.

48 hours

Tags - *optional*

Key Value

Remove

Add tag

Cancel Create training job

Model Settings

Model settings

Model name

richard-de3-prep

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

IAM role

Amazon SageMaker requires permissions to call other services on your behalf. Choose a role or let us create a role that has the **AmazonSageMakerFullAccess** IAM policy attached.

AmazonSageMaker-ExecutionRole-20230101T194235

Create role using the role creation wizard

Container definition 1

▼ Container input options

◉ Provide model artifacts and inference image location

Model Settings

richard-de-prep

Clone | Create model package | Stop | Create model

Job settings

Algorithm

Algorithm ARN
-

Training image
683313688378.dkr.ecr.us-east-1.amazonaws.com/sagemaker-xgboost:1.3-1

Input mode
File

Additional volume size (GB)
1

Maximum runtime (s)
1800

Maximum wait time for managed spot training(s)
-

Managed spot training
Disabled

Volume encryption key
-

Model Settings

▼ **Provide model artifacts and inference image options**

- (•) Use a single model
  Use this to host a single model in this container.
- ( ) Use multiple models
  Use this to host multiple models in this container.

Location of inference code image
Type the registry path where the inference code image is stored in Amazon ECR.

683313688378.dkr.ecr.us-east-1.amazonaws.com/sagemaker-xgboost:1.3-1 → **from training job**

Location of model artifacts - *optional*
Type the URL where model artifacts are stored in S3.

s3://richard-de-prep/sagemaker-autoscout/output/richard-de-prep/output/model → **from s3 (modeltargz uri)**

The path must point to a single gzip compressed tar archive (.tar.gz suffix).

Container host name - *optional*
Type the DNS host name for the container.

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

▼ **Environment variables** - *optional*

Model Settings

Network

Enable network isolation

Containers that run with network isolation can't make any outbound network calls. This is required for AWS ML Marketplace products.

VPC - *optional*

For better security, we recommend that you use a private VPC.

No VPC

Tags - *optional*

Key | Value | Remove

Add tag

Cancel | Create model

Endpoint Config Settings

Create endpoint configuration

To deploy models to Amazon SageMaker, first create an endpoint configuration. In the configuration, specify which models to deploy, and the relative traffic weighting and hardware requirements for each. SeeDeploying a Model on Amazon SageMaker Hosting Services Learn more about the API

Endpoint configuration

Endpoint configuration name

richard-de3-prep

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Type of endpoint

- (•) Provisioned
- ( ) Serverless

Encryption key - *optional*

Encrypt your data. Choose an existing KMS key or enter a key's ARN.

No Custom Encryption

Endpoint Config Settings

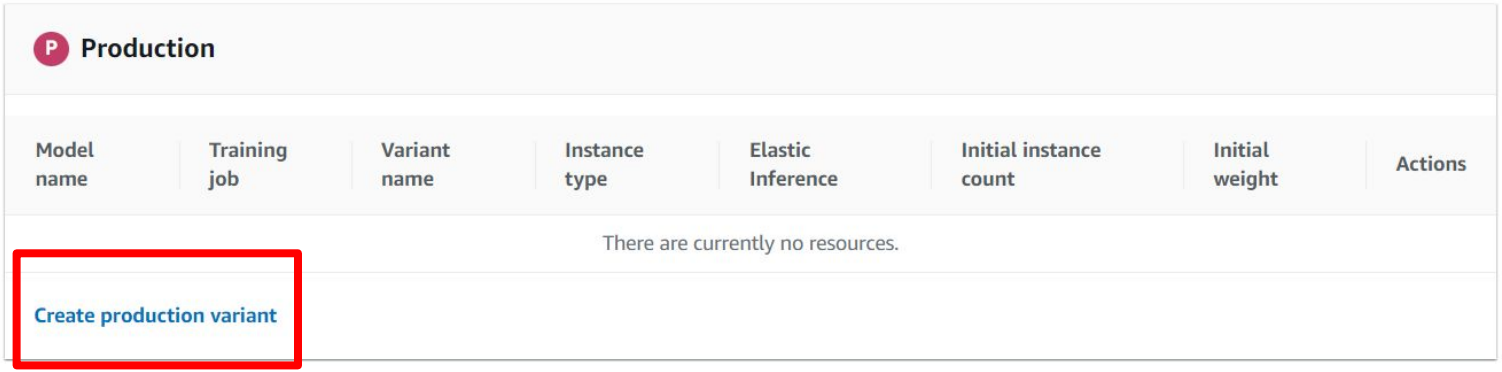

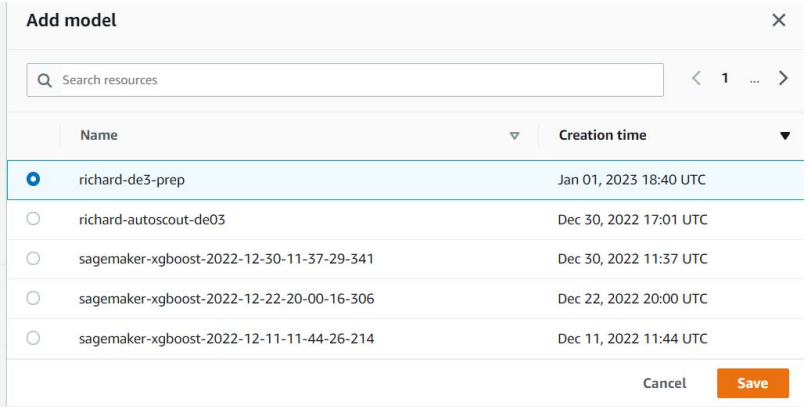

Endpoint Config Settings

P Production

Model name	Training job	Variant name	Instance type	Elastic Inference	Initial instance count	Initial weight	Actions
richard-de3-prep	richard-de-prep	variant-name-1	ml.m4.xlarge	none	1	1	Edit | Remove

Create production variant

Cancel | Create endpoint configuration

Endpoint Settings

Endpoint

Endpoint name

Your application uses this name to access this endpoint.

richard-de3-prep

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Attach endpoint configuration

- (●) **Use an existing endpoint configuration**
  Use an existing endpoint configuration or clone an endpoint configuration.
- ( ) **Create a new endpoint configuration**
  Add models and configure the instance and initial weight for each model.

Endpoint Settings

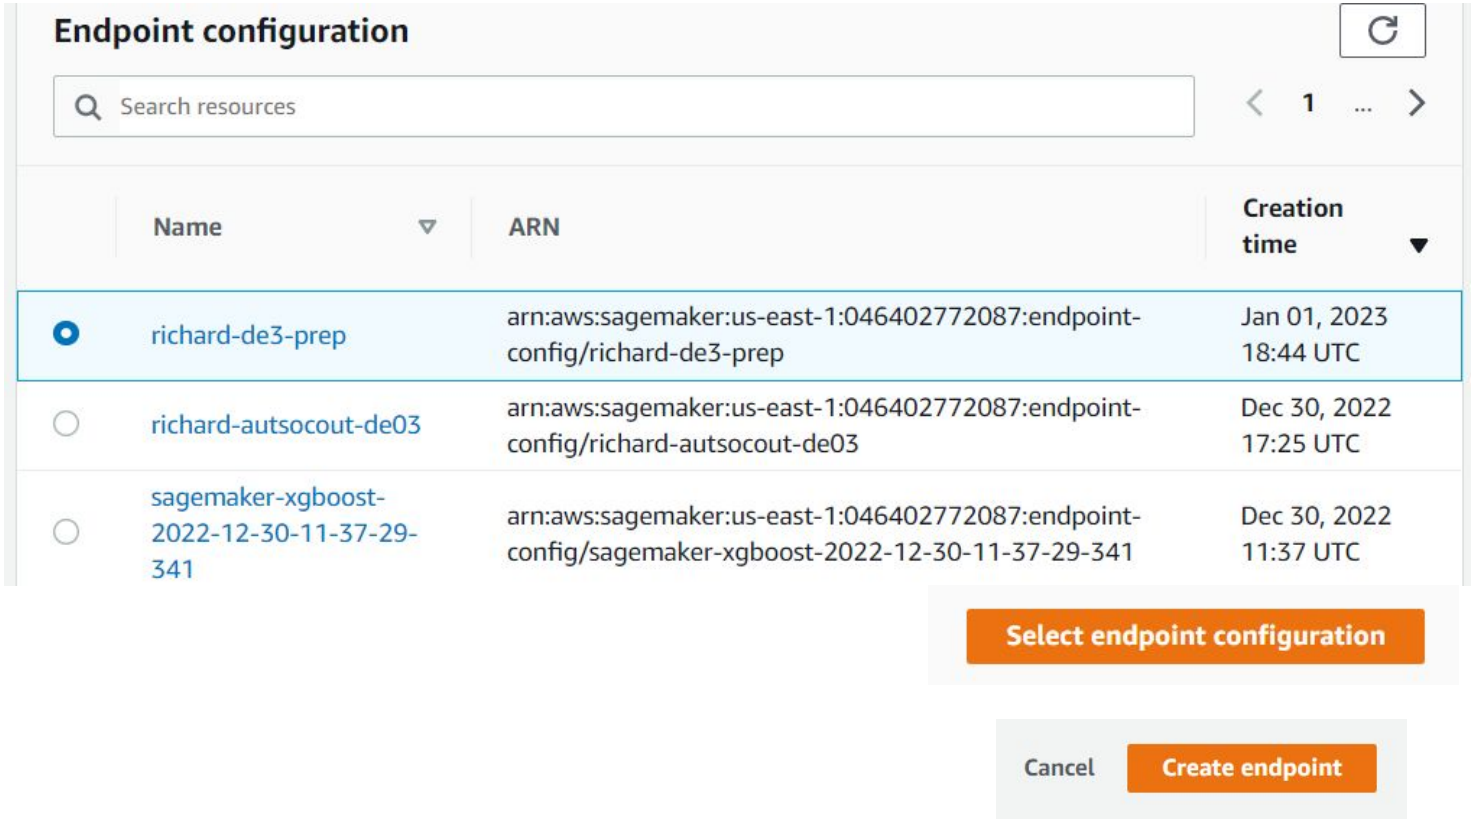

SageMaker Studio